

72-Hour Launch Execution Plan

Day 1: Foundation (Next 24 Hours)

Morning (Hours 0-4)

☐ Domain & Infrastructure

- Register gpudex.io (or gpudex.ai, gpuswap.io)
- Set up Vercel/Netlify for frontend
- Create GitHub repo
- Set up basic CI/CD

☐ Legal Foundation

- Draft basic Terms of Service
- Add disclaimer: "Prices for reference only"
- Register LLC (use Stripe Atlas)

Afternoon (Hours 5-8)

☐ Core Scraping Infrastructure

- Deploy Python backend on Railway/Render
- Implement Vast.ai scraper (they have public API)
- Add RunPod price extraction
- Set up PostgreSQL for price history

Evening (Hours 9-12)

☐ MVP Frontend

- Deploy the HTML artifact
- Connect to backend API
- Add Google Analytics
- Implement email capture

Day 2: Expansion (Hours 24-48)

Morning (Hours 13-16)

☐ Provider Integration

- Add TensorDock scraper
- Add Lambda Labs prices
- Implement Paperspace API
- Add 5 more providers

Afternoon (Hours 17-20)

☐ Advanced Features

- Build price history charts
- Add GPU availability tracker
- Implement basic arbitrage alerts

- Create provider reliability scores

Evening (Hours 21-24)

☐ **Developer Experience**

- Document public API
- Create Postman collection
- Build API key system
- Add rate limiting

Day 3: Launch (Hours 48-72)

Morning (Hours 25-28)

☐ **Polish & Optimize**

- Mobile responsive design
- Page speed optimization
- SEO meta tags
- Error handling

Afternoon (Hours 29-32)

☐ **Launch Preparation**

- Write HackerNews post
- Create ProductHunt listing
- Draft Reddit posts
- Prepare Twitter thread

Evening (Hours 33-36)

☐ **Go Live**

- Submit to HackerNews
- Post on r/LocalLLaMA, r/MachineLearning
- Launch on ProductHunt (midnight PST)
- Send to AI newsletters

Critical Success Metrics

Day 1 Targets

- ✓ Live website with real prices
- ✓ 3+ providers integrated
- ✓ 100+ unique visitors

Day 3 Targets

- ✓ 1,000+ unique visitors
- ✓ 100+ email signups
- ✓ 10+ API key requests
- ✓ 1+ partnership inquiry

Revenue Activation (Week 1)

Immediate Monetization

1. **Affiliate Links** (Day 4)
 - Vast.ai: 5% commission
 - RunPod: 10% first month
 - Apply to all programs
2. **Premium API** (Day 7)
 - \$0 - 100 calls/day
 - \$99 - 10,000 calls/day
 - \$999 - Unlimited + priority
3. **Enterprise Leads** (Day 10)
 - "Book a Demo" for bulk buyers
 - Custom routing solutions
 - SLA guarantees

Growth Hacks

Content Strategy

- "GPU Pricing Report 2025" (SEO play)
- "How I Save \$10k/Month on GPUs" (viral potential)
- Daily price alerts on Twitter

Partnership Outreach

- Email top 50 AI startups
- Offer free premium access
- Get testimonials

Technical Moat

- Start logging all price data NOW
- Build ML price prediction model
- Patent-pending on routing algorithm

The Unfair Advantages

1. **First-Mover**: No true aggregator exists
2. **Network Effects**: Prices improve with volume
3. **Data Moat**: Historical prices = prediction power
4. **Developer Love**: Clean API = viral adoption

Post-Launch Priorities

Week 2: Add 20 more providers **Week 3**: Launch mobile app **Week 4**: Raise pre-seed round **Month 2**: 10,000 MAU target **Month 3**: \$50k MRR target

Remember

You're not building a price comparison site. You're building the **infrastructure layer** for the \$50B GPU compute market. Every decision should optimize for developer adoption and transaction volume, not immediate revenue.

The market will tell you what to build next. Ship fast, listen carefully, iterate relentlessly.

Your only job for 72 hours: Get to 1,000 users who check prices daily.

"In 2025, every AI company will need GPUs. In 2026, they'll all use GPUDex to find them."